

Setanta - Developing an Irish programming language,
learning environment, and an original parser generator
for TypeScript

Eoin Davey

Final Year Project - 2019/2020 - 5 credits

B.Sc. Computational Thinking



Department of Computer Science

Maynooth University

Maynooth, Co. Kildare
Ireland

A thesis submitted in partial fulfilment of the requirements for the B.Sc. Computational
Thinking.

Supervisor: Dr. Barak A. Pearlmutter.

Contents

1	Abstract	3
2	Introduction	4
2.1	Motivation	4
2.2	Problem Statement	4
2.3	Approach	5
2.4	Evaluation Metrics	6
2.5	Project Achievements	6
3	Background	8
3.1	Executing Code on the Browser	8
3.2	TypeScript	8
3.3	PEG and CFG Grammars	9
3.4	Parser Generators	9
3.5	VSO vs SVO languages	9
4	Technical Problem Description	10
4.1	<i>tsPEG</i>	10
4.2	<i>Setanta</i>	10
4.3	try-setanta.ie	11
5	The Solution	12
5.1	<i>tsPEG</i>	12
5.1.1	High Level Design	12
5.1.2	Grammar specification	13
5.1.3	Bootstrapping	14
5.2	<i>Setanta</i>	14
5.2.1	High Level Design	14

5.2.2	Syntax	14
5.2.3	The @ operator - An example of Irish vs English host language	16
5.2.4	Semantics	17
5.2.5	Blocking operations & concurrency	17
5.3	The learning environment - try-setanta.ie	18
5.3.1	High Level Design	18
5.3.2	Graphics API	18
5.3.3	Saving Code	19
6	Evaluation	20
6.1	<i>tsPEG</i>	20
6.2	<i>Setanta</i>	20
6.3	try-setanta.ie	21
6.4	Overall end to end correctness	21
7	Conclusion	22
7.1	Goals and project reception	22
7.2	The value of non-English PLs	22
7.3	The future of the project	23

Chapter 1

Abstract

The world of programming languages is one dominated by the English language, practically all programming languages that are used today are designed to be used in English. English is established as the lingua franca of the programming world. Studies have shown that language affects the thoughts of the speaker[2]. How does it affect how we design our programming languages? In recent times Irish is often thought of as an academic, historical language. However, Irish is a language used by 73,000 people on a daily basis. There is a large contingent of fluent Irish speakers who, if they want to learn how to program, have no choice but to learn in English. This project is an exploration of the design and implementation of a programming language (*Setanta*) from the ground up, to be written in a non-English language, namely Irish, and the effects that the “host” language has on its syntax and semantics. We aim to create a novel, expressive language, and an online environment where the language can be used and learned. In this project we create a modern, powerful Irish programming language *Setanta*. In the process of designing *Setanta* we discover syntactic constructs that are motivated by the Irish language. We develop and launch an online learning platform (try-setanta.ie). When implementing an interpreter for *Setanta*, we find and fill gaps in the tooling available for creating a programming language to be executed in the browser. Specifically by creating an innovative parser generator (*tsPEG*) for the language TypeScript.

Chapter 2

Introduction

For easier distinction between programming languages and “human” languages, from this point I will refer to programming languages as **PLs**, and traditional languages as just **languages**.

2.1 Motivation

English is the language of choice for the programming world, even PLs developed in non English speaking countries are designed to be written in English, e.g., *Lua* (developed in the Netherlands), *Ruby* (developed in Japan). This focus on one single language must have some impact the way we design our PLs. Many PLs have been written for other languages, but if you go to use one you will almost certainly find that it is a *translation* of a PL originally written for English speakers[3]. If we design a PL from the ground up around a non English language, what changes do we see between it and the industry standard English PLs. By finding deep links between the English language and the PLs used in industry, I hope to reveal subtle ways in which diversity of thought is being constrained, by forcing all programming to be done through the lens of English.

Irish is a language that historically has faced significant hostility, and today finds itself a minority language in its own country, however it is still spoken by over 73,000 people daily[1]. If an Irish speaking person wishes to learn about programming, they have no choice but to do it through the medium of English. By creating an Irish PL and an online learning environment around it I hope to enable people to learn to program in they way that they want to.

2.2 Problem Statement

This project involves the design and implementation of a new, modern, innovative PL, named *Setanta*. *Setanta* is to be developed entirely in Irish. It will not be a translation or a modification of a previously existing PL, this is to allow the PL design to be influenced by the Irish language

at every stage of the design and implementation process.

Setanta will be designed with education in mind. It will be built to run in the browser, in order to enable high ease of access to as many people as possible. By running the code in the browser, no installations are required to use the PL, just a web browser.

Setanta must be a modern PL with all industry standard features, this is to ensure that by learning it, you learn the most fundamental programming concepts. Additionally it should be built to overcome the limitations of the browser environment. Executing code in the browser limits you to a single thread and no blocking operations, by building a language on top of this we can abstract out these limitations and enable users to use concurrent functions and blocking operations such as IO.

An online learning environment will be created where the user can write and *Setanta* in the browser. It should be accessible and easy to use. To assist in the learning process the environment will take inspiration from popular educational tools like *Scratch* and *Logo* and have a graphical interface where the user can draw shapes and interact with visual elements. Research has shown that the use of visual elements in educational approaches improves the learning experience[4].

To implement an interpreter for a PL a parser is needed, usually a parser generator is used to do this. However, as not many languages are built to be browser-first, the parser generator choices available for TypeScript (my PL of choice for this project) were not quite suitable. This leads to the additional part of this project to create a novel parser generator for TypeScript. The parser generator must be powerful enough to support *Setanta*, as well as to be capable of bootstrapping its own parser. It should be built on the latest innovations in parsing technology, providing accurate syntax error detection and ASTs to the user. The ASTs generated by the parser should be strongly typed, to enable maximum utility of the TypeScript type system.

2.3 Approach

The approach to completing this project involves a few main steps. First a prototype of the learning environment is made as an experimentation sandbox and proof of concept. JavaScript will be initially used as a stand in for *Setanta* as it can be already executed on the browser. After experimenting with the sandbox we move to the design process of *Setanta*. We must decide on several important features of the language. The design process will involve creating several documents outlining the decision process in real time. The linguistic properties of Irish will be contrasted with those of English and used to influence the syntax and semantics of the language.

After a design is settled on we move to creating the parser generator *tsPEG*. This is a

key component and will be needed to create the parser for *Setanta*. *tsPEG* is worked on as a largely independent project, in fact my supervisor has stated that he thinks that *tsPEG* is of sufficient independent interest to be a final year project of its own. *tsPEG* is created by using a bootstrapping process whereby a simple parser is made by hand for a very basic grammar, then this is used to self bootstrap further and further powerful features until we have created an expressive, state of the art parser generator.

tsPEG is then used to create the parser for *Setanta*, and work on implementing *Setanta* can begin. Like *tsPEG*, *Setanta* is created by first implementing the basic features of a PL, variables, loops, function calls. We then move on to implementing more and more complex features including classes, inheritance, closures, first order functions and concurrency support. As we create *Setanta* we replace the JavaScript in the learning environment with *Setanta*, thus creating the full end-to-end product. After creation of *Setanta* we return to the learning environment, adding a backend that allows users to save their code and send it to others.

2.4 Evaluation Metrics

We will evaluate how successful the project was by examining the level of success that was attained in completing the goals.

The goals for *Setanta* include creating a new, modern programming language in Irish, evaluating how it differs from languages designed in English, how beginner friendly it is, and how much interest there is in the language from the community at large. For *tsPEG* the goals are to create a powerful parser generator for TypeScript. *tsPEG* should be powerful enough to build a real programming language with, it should be capable of self-hosting its own input parser. *tsPEG* should be based on novel ideas in parser technology, and be highly configurable to different projects. The learning environment must be an accessible website where users can write and execute *Setanta* code. The website should allow users to experiment with *Setanta*, both through classical text IO programming, and by allowing them access to an API to draw and manipulate shapes and objects on a graphics display.

2.5 Project Achievements

In the completion of this project several significant accomplishments were realized.

- **Successful creation and launch of *tsPEG***

The TypeScript parser generator *tsPEG* was created and released on the NPM (Node Package Manager). *tsPEG* was announced on some TypeScript forums and garnered some

attention, peaking at 242 average weekly downloads for a period after its release.

- **Implementation of *Setanta*.**

As the main goal of the project, *Setanta* was developed almost exactly to the original design. It's a modern, powerful, expressive language with all the bells and whistles that we have come to expect. *Setanta* is a strong but dynamically typed language that can be executed both in the browser and locally on the command line through the NodeJS runtime.

- **try-setanta.ie - The *Setanta* learning environment.**

try-setanta.ie satisfies all the requirements that were expected of it. Visiting try-setanta.ie the user can write and execute *Setanta* code. They are able to save their code and send it to others with unique URLs assigned to each saved piece of code. try-setanta.ie also features a graphics display where the programmer can draw and manipulate different shapes and objects.

- **Allowing Irish to affect the design of *Setanta*.**

Part of the goals of designing *Setanta* was to allow the Irish language to influence its design, and by doing so to see what parts of industry standard languages are subtly influenced by their design in English. *Setanta* has some syntactic constructs that were affected by the use of Irish as its “host” language. These effects highlight some of the Anglo-centric designs in today's languages, including in features that we consider to be almost universal.

- **Abstraction out over the browsers concurrency issues.**

Anyone who's written code for the browser will know the pain of the single-threaded, non-blocking world of JavaScript. Programs cannot launch multiple threads, or wait for operations to finish. The *Setanta* runtime abstracts out the JavaScript internal task queue and callback systems, and allows the user to write code that waits for blocking operations such as user input, and to execute things concurrently such as running infinite rendering loops but still allowing keyboard inputs to be processed.

Chapter 3

Background

3.1 Executing Code on the Browser

The only PL that standard web browsers (*Google Chrome, Mozilla Firefox, Safari, Edge, Internet Explorer, etc.*) can execute is JavaScript. There is a movement to bring other PLs to the browser via a technology called WebAssembly, but this has not fully come to fruition yet.

In addition to this restriction on PLs, each browser instance limits executing code to a single, non-blocking thread. To overcome the non-blocking limitation of the execution thread the JavaScript environment exposes an API to a simple task queue, allowing users to pass in **callbacks**, which are functions that will be executed at some later time. In Figure 3.1 we compare code in a version of JS that allows blocking, and show the real code that must be used using callbacks using the `setTimeout` API to add the callback to the task queue.

<pre>1 // Would work in an ideal world 2 function sleepFor100ms() { 3 print('Starting sleep') 4 sleep(100) 5 print('After sleep') 6 }</pre>	<pre>1 // The real way to achieve this 2 function sleepFor100ms() { 3 print('Starting sleep') 4 setTimeout(() => print('After sleep'), 100) 5 }</pre>
---	--

Figure 3.1: Code if thread can be blocked vs real situation.

3.2 TypeScript

TypeScript is an open source PL being developed by Microsoft, TypeScript is the primary PL that is used in the implementation of this project. Microsoft describes it as

TypeScript is a typed superset of JavaScript that compiles to plain JavaScript.[6]

TypeScript allows developers to write JavaScript programs with all the benefits of compile time static type checking.

3.3 PEG and CFG Grammars

Formal grammars are sets of rules for re-writing strings, first formally researched and defined by Noam Chomsky, 1956[9]. In recent times they find frequent use in specifying the syntactic grammar of various PLs.

The formal grammar family that has found the most usage is that of CFGs, Context Free Grammars. PEG Grammars are a new type of grammar that has come into prevalent usage recently, they are very similar to CFG grammars but are defined in such a way that they cannot be ambiguous. It is conjectured that PEGs are more powerful than CFGs[10]

3.4 Parser Generators

A parser generator is a program that, given a description of a grammar in some formal syntax, outputs a parser for that grammar. Parser generators have existed since the earliest days of computer programming. One of the earliest major parser generators, YACC, was released in 1975 [5], and was based on LR parsing.

Several parser generators are available for JavaScript, including the excellent *pegjs* generator. However, no such parser generator was available for TypeScript, this is what prompted the creation of *tsPEG*.

3.5 VSO vs SVO languages

In linguistics the order of how the different categories of words appear in sentences is often studied. One order that is studied is called the “constituent word order”. This is the order that a verb (V), object (O), and subject (S) appear in a sentence[11]. English is an SVO language, meaning that the subject comes first, then the verb, then the object. Irish is a VSO language, meaning the verb comes first, then the subject, then the object.

Chapter 4

Technical Problem Description

4.1 *tsPEG*

tsPEG must be designed to meet the following requirements:

- *tsPEG* must be a fully functional TypeScript parser generator. It should allow the user to specify a PEG grammar in some input format, and output a fully correct parser for that grammar.
- *tsPEG* should also be written in TypeScript, to allow it to self-host, and bootstrap. Hence *tsPEG* will run on the NodeJS JavaScript runtime.
- The outputted parser must correctly match precisely the input grammar, and report any syntax errors encountered along the way.
- The parser must utilise the TypeScript type system to the fullest extent, making use of discriminated union types and enum types to create a strongly typed AST for the input grammar.
- The *tsPEG* input format should be expressive enough to allow users to specify complex and powerful grammars easily. It should support the standard set of PEG operators.

4.2 *Setanta*

The technical requirements of *Setanta* are:

- *Setanta* should be fully usable with no knowledge of English, only Irish can be relied on.
- *Setanta* is required to be executable in the browser as well as locally.

- The syntax and semantics of *Setanta* must be directly influenced by the linguistic and cultural properties of the Irish language, a simple re-skin of an existing language with Irish keywords is not sufficient.
- Standard features of modern PLs must be present. The domain of *Setanta* is education, it must not be esoteric, learning *Setanta* should teach fundamental programming concepts.
- When the opportunity presents itself *Setanta* should make programming in it as accessible as possible. This includes using simplified vocabulary compared to mainstream languages, as well as not requiring the use of fadas (diacritics like áéíóú), as not all users can type these characters.

Outside of these technical requirements, the design of *Setanta* is largely free for me to decide.

4.3 [try-setanta.ie](#)

The learning environment at [try-setanta.ie](#) also has some technical requirements. [try-setanta.ie](#) exists as a portal to using and learning *Setanta*. Therefore it must satisfy the following conditions:

- [try-setanta.ie](#) must enable visitors to the site to write and execute *Setanta* code.
- [try-setanta.ie](#) must be clearly laid out and simple to use.
- Some method of saving code and sharing it with friends needs to be included.
- As with *Setanta*, [try-setanta.ie](#) must not require any knowledge of English, only Irish can be used.
- Some sort of graphical display must be included, and should have an API that is exposed to the *Setanta* runtime.

Chapter 5

The Solution

I will discuss the design and implementation of each of the main components of the project separately.

5.1 *tsPEG*

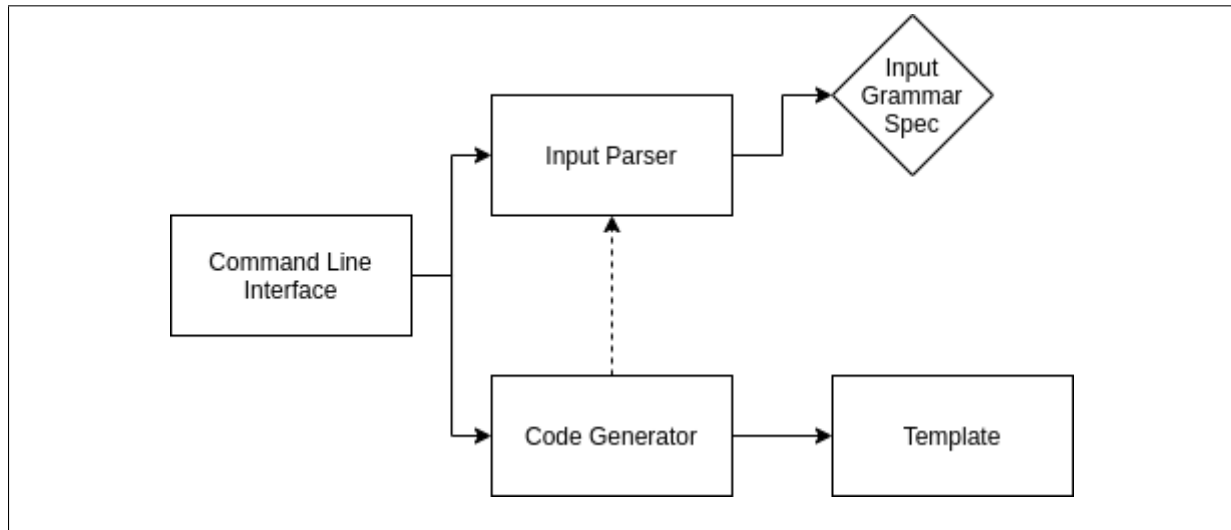
5.1.1 High Level Design

tsPEG was built using the TypeScript PL. *tsPEG* runs on the NodeJS JavaScript runtime. The NodeJS runtime allows us to execute JavaScript (and hence TypeScript) code locally.

The source code is hosted online at github.com/EoinDavey/tsPEG, and the *tsPEG* package is available on the NPM package repository at npmjs.com/package/tspeg.

tsPEG is self hosting, meaning that the input parser for *tsPEG* was generated by *tsPEG*. The software architecture of *tsPEG* is given in Figure 5.1. The usage flow of using *tsPEG* is as follows. First the user creates a grammar file, specifying the input grammar they want to generate a parser for. The syntax for the grammar file follows a similar pattern to the familiar EBNF syntax for grammar specification. In this file they specify the syntax rules, as well as names of AST fields, and computed properties. The *tsPEG* binary is then called and it is passed in the grammar file.

*tsPEG*s own parser consumes the grammar file, and generates an internal description of the grammar. The code generator then takes this grammar specification and generates parsing functions for each rule, as well as TypeScript type declarations and classes for each of the AST nodes. These functions and types are then packaged up into a template file and then written to disk.

Figure 5.1: *tsPEG* High Level Design

```

1 SUM := head=FAC tail={ op='\+|- ' sm=FAC }*
2 FAC := head=ATOM tail={ op='\*/' sm=ATOM }*
3 ATOM := val=INT
4       | '\(' val=SUM '\)'
5 INT  := val='[0-9]+'
  
```

Figure 5.2: *tsPEG* input grammar example

5.1.2 Grammar specification

tsPEG uses a custom syntax to define grammars. Figure 5.2 contains an example of a grammar specification for simple arithmetic expressions like "1+2*3".

Grammars are defined by a sequence of grammar rules, for example

$$\text{match} := \text{rule1} \mid \text{rule2} \mid \text{'a+'}$$

defines a new rule **match** that tries first **rule1** then **rule2**, then tries to match the regex expression **a+**. The full list of operators is included in Appendix A.

The *tsPEG* grammar also allows specification of computed properties, for example Figure 5.3 defines a rule to match integer literals that stores the value of the integer as a computed property.

```

1 INT := literal='[0-9]+'
2      .value = number { return parseInt(this.literal) }
  
```

Figure 5.3: *tsPEG* computed properties example

5.1.3 Bootstrapping

When developing *tsPEG*, first a simple input parser was written by hand, supporting only the most basic syntax. A code generator was written that could take in the AST from this simple grammar and output a parser for it. Then the meta-grammar for the input grammar syntax was written and the generator was run. This replaced the hand written parser with a new generated one, self hosting itself and opening itself up to bootstrapping.

This new self hosting generator was then used to add more and more features and operators to itself, eventually resulting in the full self describing meta-grammar for the *tsPEG* input syntax included in Appendix 1.

5.2 *Setanta*

5.2.1 High Level Design

Setanta is also written in TypeScript. The code structure of *Setanta* allows it to be run as a local REPL through the NodeJS engine, but also in the browser. A high level architecture diagram is given in Figure 5.4.

The *Setanta* runtime uses the *tsPEG* parser generator to generate its parser. Both the browser and CLI execution environments import the same interpreter and parser, but provide their own set of built-in functions. Each runtime environment passes some input text to the parser, the parser returns an AST or a syntax error object. This AST can then be passed to the interpreter with the relevant builtins. The interpreter can then be stopped and started asynchronously from the main executing thread.

5.2.2 Syntax


The syntax of *Setanta* is new, but should feel familiar to most people. It has been designed to simple and approachable.

Setanta programs, like most imperative languages, consist of a sequence of statements. Some important *Setanta* features are outlined below:

- **Variable declarations**

In *Setanta* variables are declared using the `:=` operator, and can be re-assigned using the classic `=` operator. The distinction is to provide a clear lexical difference between variable declaration and reassignment.

- **Loops + Conditionals**

Figure 5.4: *Setanta* High Level Design

Setanta support the classic conditional execution structure of if, else if, else. This is mostly a direct translation into Irish. However it should be noted that no bracketing is required around the expression.

Listing 5.1: *Setanta* conditionals

```

1  má x == 0
2      scríobh('Tá x cothrom le 0')
3  nó má x == 1
4      scríobh('Tá x cothrom le 1')
5  nó
6      scríobh('Tá x níos mo ná 1')
```

Setanta supports two main types of loops, “le idir” loops that allow the user to specify start and ends to the loop, and “nuair-a” loops, which are the familiar while loops.

Listing 5.2: *Setanta* loops

```

1  i := 0
2  le i idir (0, 10)
3      i = i + 1
4  x := 0
5  nuair-a x < 10
6      x = x + 1
```

- **Classes + Functions**

Setanta supports declaring new classes, with methods, as well as functions on their own.

Listing 5.3: Setanta classes

```

1 creatlach Person ó Animal {
2     gníomh nua(name) {
3         name@seo = name
4     }
5     gníomh speak() {
6         scríobh('Hi, My name is ' + name@seo)
7     }
8 }
```

• Literals

Setanta supports literals for integers, booleans, null, strings and lists.

Listing 5.4: Setanta literals

```

1 a := 500
2 b := 'Dia duit domhan'
3 c := [1,2,3,4, fíor ]
4 d := fíor != breag
5 c := neamhní
```

5.2.3 The @ operator - An example of Irish vs English host language

One of the most important features of *Setanta*'s syntax is the @ operator. The @ operator is where the influence of the Irish language on the design of *Setanta* is felt most clearly. In fact, the @ operator is a simple innovation that addresses 2 distinct differences between Irish and English.

The @ operator is the lookup operator, it is used to reference member fields and methods of objects. It is functionally equivalent to the classic dot ('.') operator in C, C++, Java, Python, JavaScript etc. The subtle difference is the order, both the @ operator and the '.' operator are binary operators, but the @ operator flips its arguments compared to the '.'.

In the standard usage of the dot operator, the expression `a.b` refers to the member “b” of the object “a”. In *Setanta* `a@b` refers to the member “a” of the object “b”, the order has been reversed. This change seems unimportant, however, it breaks 2 deep and subtle links between English and the standard way we look at OOP programming languages.

Firstly, as discussed in the technical background (3.5), English is an SVO language. This property of English is directly reflected in the design of the classic dot operator.

The usage of the dot operator for member lookup gives rise to expressions like `man.goTo(shop)`, or `dbClient.query(q)`, these expressions implicitly put the subject first, then the verb, then the object, directly mimicking the SVO structure of the host language.

By using the `@` operator in *Setanta* we instead get expressions like `goTo@man(shop)` or `query@dbClient(q)`, these expressions put the verb first, then the subject, then the object, reflecting the VSO structure of Irish. This subtle connection between a simple operator like `'.'` and the linguistic properties of English is not immediately apparent.

There is a second connection between English and the lookup operation that is broken by the `@` operator. In English when talking about possession relationships (“has-a” relationships in database terms), we list the entities in decreasing order of possession, e.g.,

The man’s car door window pane

The order of possession here is

$$man \rightarrow car \rightarrow door \rightarrow window \rightarrow pane$$

In direct contrast to this, in Irish we list possession in increasing order, meaning we start at the bottom of the possession relationship and work our way up. We would instead list

$$pane \leftarrow window \leftarrow door \leftarrow car \leftarrow man$$

This linguistic difference between Irish and English is again addressed by the `@` operator. The standard `'.'` operator reflects the English language structure, `man.car.door.window.pane`, but in *Setanta* we use the Irish ordering `pane@window@door@car@man`. This is another connection between English and the standard design of PLs that is easily overlooked.

5.2.4 Semantics

The semantics of *Setanta* will be familiar to most users, it’s an imperative, strongly dynamically typed language. *Setanta* supports a gauntlet of modern features including first class functions, inheritance, event based concurrency, and automatic memory management. Although *Setanta* is an imperative language, it does support a more functional style of programming as well.

5.2.5 Blocking operations & concurrency

As discussed in the technical background section, the JavaScript runtime, whether it be NodeJS or a browser, only allow usage of a single thread, and no blocking operations. However, writing *Setanta* allows the user to use blocking functions. Figure 5.5 shows a side by side comparison

Listing 5.5: JavaScript	Listing 5.6: Setanta
<pre> 1 setTimeout(() => { 2 console.log('sleep1'); 3 setTimeout(() => { 4 console.log('sleep2'); 5 setTimeout(() => { 6 console.log('sleep3') 7 }, 100); 8 }, 100); 9 }, 100); </pre>	<pre> 1 coladh(100) 2 scríobh('coladh1') 3 coladh(100) 4 scríobh('coladh2') 5 coladh(100) 6 scríobh('coladh3') </pre>

Figure 5.5: Equivalent code in JavaScript + Setanta

of equivalent programs in JavaScript and Setanta, it's clear that the Setanta program is much simpler due to allowing the program to block.

If you open the JavaScript console on any browser and enter `while(true){}` the browser tab will cease to be functional, i.e no buttons or text boxes or animations will work anymore. This is because the JavaScript engine is single threaded. In contrast, if you write `nuair-a fíor {}` in *Setanta*, the browser tab will remain responsive. This is due to the implicit concurrency of all *Setanta* functions.

5.3 The learning environment - try-setanta.ie

5.3.1 High Level Design

try-setanta.ie, like the other components is written mainly in TypeScript. The website is built on Google's LitElement library. LitElement is a library to write custom re-usable web components, it's a descendant of the Polymer framework, which has since been deprecated. LitElement is used to achieve a Material Design aesthetic. The text editor component uses the open source library CodeMirror. The backend for try-setanta.ie is written in Go, and runs on the Google App Engine (GAE) cloud infrastructure platform. The backend for the site supports the ability to save and retrieve *Setanta* programs to the cloud.

5.3.2 Graphics API

To better enable the educational experience on try-setanta.ie, the *Setanta* runtime has access to a graphics API. The graphics API allows the user to draw simple graphics primitives and animations. This API is exposed to the *Setanta* program as a global object. The user can access various fields and methods on this object to draw and manipulate shapes on the display. Figure

5.6 shows a short program on try-setanta.ie that draws a Sierpinski triangle. The graphics API also exposes methods to consume keyboard inputs, so the user can program simple games like *Snake* with ease. It also features a full-screen mode.

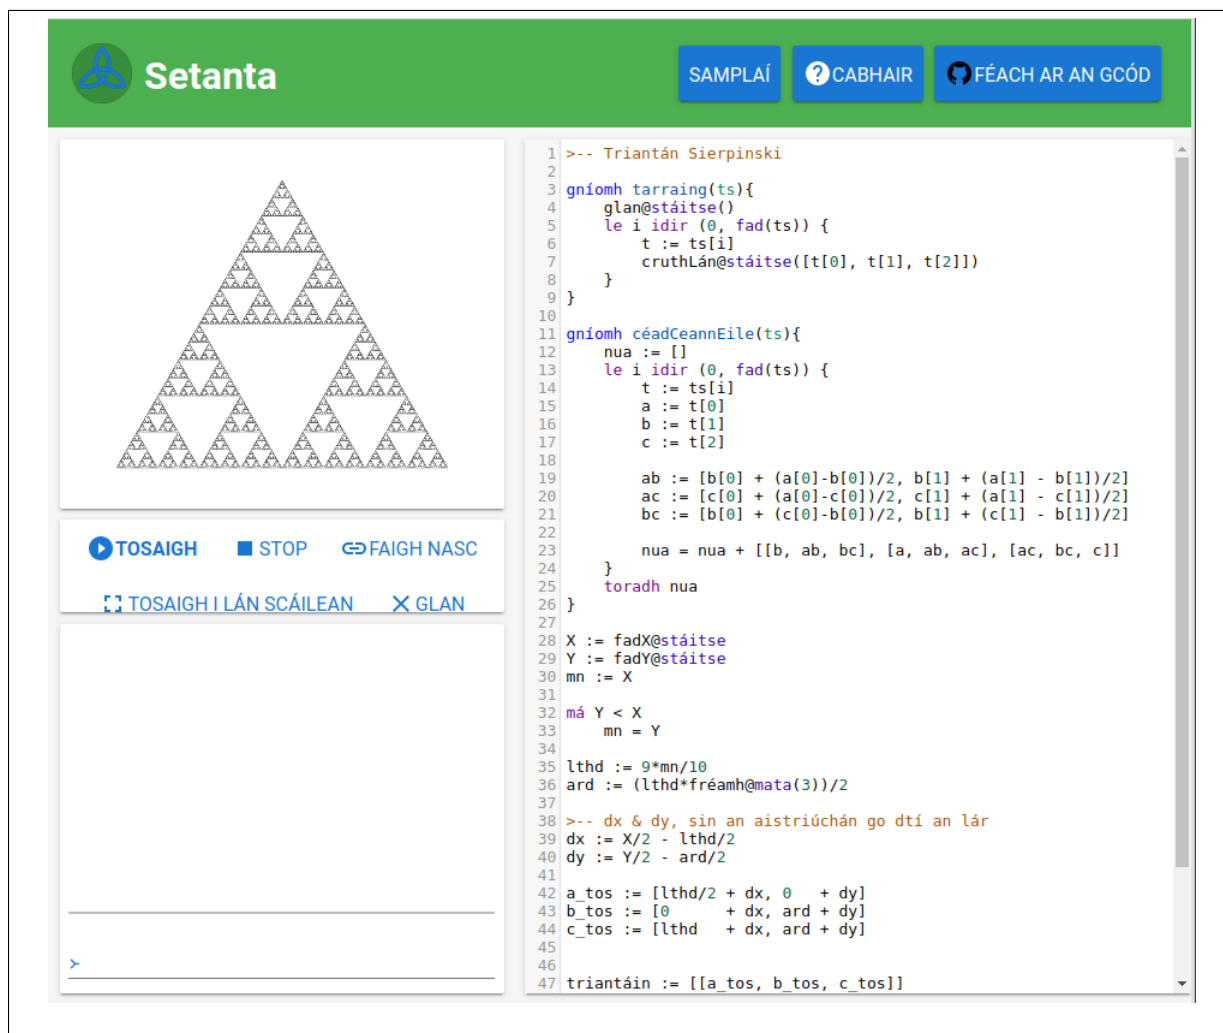


Figure 5.6: Sierpinski Triangle drawn with *Setanta*

5.3.3 Saving Code

try-setanta.ie supports the ability to save a program to a specific URL. Clicking “Faigh nasc” (meaning “get link”) redirects the user to a unique URL with which they can access their code. For example, the Sierpinski triangle example is available at:

try-setanta.ie/EhEKBINjcmlwdBCAgIDgyeODCg.

This feature is handled by the Go backend on the Google App Engine. The programs are stored on the cloud Datastore, each program is identified by a unique integer. This unique integer is encoded into a string key, and this key is the URL where the code can be retrieved.

Chapter 6

Evaluation

6.1 *tsPEG*

The *tsPEG* parser generator was implemented almost precisely to the original design. The correctness and effectiveness of *tsPEG* can be seen in two direct ways.

- *tsPEG* is self-hosting, meaning that the input parser for *tsPEG* is generated by *tsPEG* itself. This would not be possible if *tsPEG*s parsers were not correct. The input grammar makes use of all the main features of *tsPEG*, so this acts as a rigorous self-referential test, whereby the whole project fails to even compile if it is not correct.
- *tsPEG* generates the parser for the *Setanta* interpreter. This proves that *tsPEG* is capable of generating parsers for fully featured programming languages, with complex and intricate grammars. The use of *tsPEG* in the *Setanta* build process serves as a full stress test of all the features available to it.

tsPEG has also attracted attention from the TypeScript community, showing that it serves a useful purpose as a tool for the community. However *tsPEG* still has many features that it would benefit hugely from, the main one being improving the syntax error reporting capabilities, and efficiency.

6.2 *Setanta*

As outlined in the solution section of *Setanta* (5.2), the implementation of *Setanta* was a success, it meets every goal that was set out originally for it.

The correctness of *Setanta* is ensured by a large suite of tests, including unit tests, and full end to end tests. In fact *Setanta* contains over 1300 lines of test code alone. These tests

ensure that every edge of language behaviour is correct, these tests are run automatically by the continuous integration system with each commit. The current state of *Setanta*'s builds can be seen at travis-ci.com/EoinDavey/Setanta.

Setanta's main issue is its speed. As it is built on top of the JavaScript environment, but requires blocking operations and concurrency, it means that *Setanta* has taken on an extremely large overhead. The result of this is that *Setanta* is notably slow at some tasks. However, the domain of *Setanta* is education, not high performance or production computing, so this is not a direct detractor from the quality of the project.

The fact that the same *Setanta* interpreter can be used in the browser and in a local CLI shows the code is well abstracted and isolated, the main Interpreter class is very portable and can be loaded in any number of environments.

Setanta has received a notable amount of attention online, before I had even announced the project I had received requests to use the project from people who had stumbled upon the repository. I gave a talk on *Setanta* at SISTEM 2020 (A tech conference held in UCD) and it received high praise from many of the attendees.

6.3 try-setanta.ie

The try-setanta.ie learning environment has also been largely successful, and meets most of the original design goals for it. try-setanta.ie is a simple to use website, where the user can write and execute *Setanta* code, it has a good editor with syntax highlighting and a powerful graphics API. The implementation of try-setanta.ie is quite modern, taking advantage of the latest features of HTML5 and CSS3. try-setanta.ie is built on the web components technology, this means that the main components of the try-setanta.ie website are fully isolated, abstracted and re-usable.

try-setanta.ie doesn't meet the original design plan of having an extremely high level graphics API, as I found in the development process of this API that it was simpler to understand a simple API that let you call a function to draw circles, squares etc. than a very high level API where you would need to use classes and inheritance and other features to interact with the API.

6.4 Overall end to end correctness

The act of going to try-setanta.ie, writing a game in *Setanta*, executing it, and playing the game in the browser is a full end to end proof that each component, try-setanta.ie, *Setanta* and *tsPEG* are all fully functional, as each part is 100% vital to the project as a whole.

Chapter 7

Conclusion

7.1 Goals and project reception

The main goals for the project were to create an original Irish programming language, a learning environment, and a parser generator for TypeScript. In reference to these goals, the project has been successful. All of the components of the project have been completed and released to the world.

At the time of writing try-setanta.ie has been visited over 400 times, and the *Setanta* REPL has been downloaded over 100 times, these parts of the project haven't even been properly announced yet, only to people I directly have spoken to.

tsPEG has been downloaded well over 700 times, and has been maintaining an average of over 20 downloads a week.

7.2 The value of non-English PLs

In creating *Setanta*, we intended for the design process to expose some links between standard PLs and English. As discussed in the section on the @ operator (5.2.3), even simple, extremely common features such as the dot (.) operator have subtle ties to the English language. Although the links we found were not groundbreaking by any means, the existence of these links is important. Diversity of thought is hampered when all programming is passed through the lens of the English language, and in this project we've seen that it's not just the keywords that form this lens, but the deeper design of the PL too.

Setanta wasn't created with the goal of being a huge technical breakthrough, in fact, as an educational language, we were aiming for familiarity. *Setanta* was created to break the pattern of English PLs, and reveal the limitations that had been placed upon us by them.

While I was working on this project, a student in Carnegie Mellon had a similar idea, and created the worlds first classical Chinese PL[12]. This shows that there is a general interest in non-English PLs in the world.

7.3 The future of the project

This project has been a passion project for me, and I fully intend to continue its development. I plan on a full release of version 1.0 of *Setanta* soon, after I complete a full documentation and tutorial for it. I hope that *Setanta* finds use in CS education, there are many schools and institutions around the country that I hope will have some interest in it. I have met many people who have shown great enthusiasm for the project, I plan on opening the project up to the open source community, hopefully including someone with much better Irish than me.

Bibliography

- [1] CSO, *The Irish language*, 2017,
cso.ie/en/media/csoie/releasespublications/documents/population/2017/7._The_Irish_language.pdf
- [2] *Language regions of brain are operative in color perception*, Wai Ting Siok, Paul Kay, William S. Y. Wang, Alice H. D. Chan, Lin Chen, Kang-Kwong Luke, Li Hai Tan, Proceedings of the National Academy of Sciences May 2009, 106 (20) 8140-8145; DOI: 10.1073/pnas.0903627106
- [3] Non-English based programming languages
https://en.wikipedia.org/wiki/Non-English-based_programming_languages.
- [4] Graphic Organisers: A review of Scientifically Based Research, The Insitute for the Advance-ment of Research in Education at AEL
<http://www.inspiration.com/sites/default/files/documents/Detailed-Summary.pdf>
- [5] Johnson, Stephen C. (1975). Yacc: Yet Another Compiler-Compiler
- [6] <https://www.typescriptlang.org/>
- [7] ECMAScript® 2017 Language Specification (ECMA-262, 8th edition, June 2017),
<https://www.ecma-international.org/ecma-262/8.0/index.html>
- [8] Mozilla Developer Network Documentation, Promises,
https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/Promise
- [9] Chomsky, Noam (Sep 1956). “Three models for the description of language”. IRE Transactions on Information Theory. 2 (3): 113–124. doi:10.1109/TIT.1956.1056813.
- [10] Ford, Bryan p (2004). “Parsing Expression Grammars: A Recognition Based Syntactic Foun-dation”. Proceedings of the 31st ACM SIGPLAN-SIGACT Symposium on Principles of Pro-gramming Languages. ACM. doi:10.1145/964001.964011. ISBN 1-58113-729-X.
- [11] Song, Jae Jung (2012), Word Order. Cambridge: Cambridge University Press. ISBN 978-0-521-87214-0 & ISBN 978-0-521-69312-

[12] <https://spectrum.ieee.org/tech-talk/computing/software/classical-chinese>